

## RL Lecture-5 Model free control

self

On-policy  $\rightarrow$  learning on the job, observing <sup>self</sup> behaviour

Off-policy  $\rightarrow$  learning following someone else's behaviour.

~~Model~~

$\rightarrow$  optimize the value function of an unknown MDP.

Why?

$\rightarrow$  most common problems are either:

i) model / MDP unknown, but experience can be sampled


ii) MDP known, but it is too expensive to compute even 1 step using the full model, except by sampling

On-policy learning  $\rightarrow$  learning from <sup>on</sup> job  $\rightarrow$  while following policy  $\pi$ , learn that policy  $\pi$ .

Off-policy learning  $\rightarrow$  observe someone do something and sampling trajectories from human's behavior (eg.), learn self behaviour (eg. robot).

In DP, alternating between policy evaluation & iteration converges to the optimal policy

We only used iterative methods, we'll try to change ① & ②.



Eg. Monte Carlo control:

1-way  $\rightarrow$  do Monte Carlo policy evaluation & then do greedy policy improvement.

Problems:

① computing  $V(s)$  requires a model, but we want model free

$$V(s) = \max_a \left( R_s^a + \gamma \sum_{s'} P_{ss'}^a V(s') \right)$$

But in model free, <sup>we</sup> want to learn a policy without any model. The Q-function defined earlier  $\rightarrow$  means that final score we would get at the end of a game (eg.) if we took action  $a$ , when in state  $s$ .  $\rightarrow Q(s, a)$ .  
So, no model dependence now. Just a state and possible actions that can be taken.

$\rightarrow$  greedy over  $Q(s, a)$   $\nearrow$  action value function

$$\pi(s) = \underset{a}{\operatorname{argmax}} Q(s, a)$$

For each  $(s, a)$  pair taken across all experience, to compute  $Q(s, a)$

- Proposal changes:

$\rightarrow$  MC evaluates  $Q = Q_{\pi}$

$\rightarrow$  Greedy policy improvement?

2nd problem  $\rightarrow$  choosing greedily  $\rightarrow$  we never visit some states or take some actions which means we ~~are~~ don't evaluate them correctly & hence, not selecting them.

eg. to illustrate that point.  $\rightarrow$  suppose 2 ~~states~~ <sup>states</sup>  $\rightarrow s_1, s_2$

upon choosing both once, we found:

$$V(s_1) = -1, \quad V(s_2) = +1 \quad \text{we choose } s_2 \text{ greedily}$$

now say,  $V(s_2) = +3 \rightarrow$  mean  $= 2$ , again choose  $s_2$  and so on...

This shows that we haven't even considered state  $s_1$  only after 1 seeing and hence, don't know what value ~~it~~ <sup>it</sup> might have as we keep choosing  $s_2$  infinitely.

$\epsilon$ -greedy exploration  $\nwarrow$  solution to ensure continual exploration

$\rightarrow$  prob.  $\epsilon \rightarrow$  choose random action

$\rightarrow 1 - \epsilon \rightarrow$  choose the greedy option.



Thus:  $\pi(a/s) = \begin{cases} \epsilon/m + (1-\epsilon) & \rightarrow \text{for greedy action} \\ \epsilon/m & \rightarrow \text{for any other action} \end{cases}$

→ keep exploring

$\epsilon$  greedy → policy improvement (ensures)

$\pi, \pi' \rightarrow$  both  $\epsilon$ -greedy policies

$$V_{\pi'}(s) \geq V_{\pi}(s).$$

$$V_{\pi}(s, \pi'(s)) = \sum_a \pi'(a/s) q_{\pi}(s, a)$$

$$= (1-\epsilon) \max_a q_{\pi}(s, a) + \frac{\epsilon}{m} \sum_a q_{\pi}(s, a)$$

Now, max,

greater than  
any weighted  
sum

$$\geq (1-\epsilon) \sum_a \frac{\pi(s/a) - (\epsilon/m)}{1-\epsilon} q_{\pi}(s, a) + \frac{\epsilon}{m} \sum_a q_{\pi}(s, a)$$

$$= \sum_a \pi(s/a) q_{\pi}(s, a) = V_{\pi}(s)$$

$\therefore V_{\pi'}(s) \geq V_{\pi}(s)$ . (From earlier similar derivations)

$\therefore$  Finally:

- ① MCPE (MCPE) policy evaluate  $Q = q_{\pi}$  ← softening of the greedy policy
- ②  $\epsilon$ -greedy policy improvement

Idea → no need to completely ~~wait~~ evaluate the policy, when we can get a better policy in only a few initial steps (similar to DP lecture).

One extreme → do update every episode, basically on the most fresh ~~data~~ information/ estimate the value function.

after every episode

① MCPE →  $Q = q_{\pi}$

②  $\epsilon$ -greedy → For those states

↑ which were visited in the episode.

How to ensure we get  $\pi^*$ ?

→ Trade off betw exploration & exploitation.

as when at  $\pi^*$ , we won't have much exploration.

Idea for balancing the 2 ideas  $\rightarrow$  GLIE.

$\rightarrow$  come up with schedule such that 2 conditions are met:

① all state-action pairs explored infinitely many times

$$\lim_{k \rightarrow \infty} N_k(s, a) \rightarrow \infty \quad \text{greedy}$$

② converges to the optimal policy

$$\lim_{k \rightarrow \infty} \pi_k(a|s) = I(a^* = \arg \max_{a' \in A} Q_k(s, a'))$$

One idea  $\rightarrow$  schedule / decay  $\epsilon$  for  $\epsilon$ -greedy.

$\rightarrow$  hyperbolic schedule  $\rightarrow \epsilon_k \approx 1/k$ .

### GLIE MC control

$\rightarrow$  sample  $k^{\text{th}}$  episode following  $\pi$

$$\downarrow \{s_1, A_1, R_1, \dots, R_T\}$$

$\rightarrow$  update mean  $Q$  value &  $N(s_t, A_t)$

$$N(s_t, A_t) = N(s_t, A_t) + 1$$

$$Q(s_t, A_t) = Q(s_t, A_t) + \frac{1}{N(s_t, A_t)} (G_T - Q(s_t, A_t))$$

$\downarrow$   
this mean is not the actual statistical mean as we are accumulating values as we improve the policy

$\downarrow$   
changing over time

$\rightarrow$  Improve policy

$$\epsilon \rightarrow 1/k$$

$$\pi \rightarrow \epsilon\text{-greedy}(Q)$$

GLIE makes sure that these collected statistics converge to the actual ~~mean~~ value over time.



$$Q(s, a) \rightarrow Q^*(s, a)$$

Iterate over the entire process.

→ considerably more efficiently than updating after a batch

Initialization → In this case, for the 1<sup>st</sup> time we observe

a state  $N(s_t, a_t) = 1$ , and  $Q(s_t, a_t) = G_T$ , so init doesn't matter. But for weighing other than  $\frac{1}{N(s_t, a_t)}$ , it will affect much more.

Now, we'll show TD.

TD v, MC


⑦ variance low → online → incomplete sequences.

- look at  $Q(s, a)$  for model-free
  - slot in TD learning for policy evaluation
  - $\epsilon$ -greedy
- } can update after 1 step., no need to wait till episode end.

so, update policy after each step.

called SARSA.

sample from environment



$$Q(s, A) \leftarrow Q(s, A) + \alpha (R + Q(s', A') - Q(s, A))$$

SARSA update.

In every time step, update value function.

→ started with  $s$ , took action  $A$  (from current policy) and reached  $s'$  → update only  $Q(s, A)$  → only for this pair

# SARSA algo for on-policy

- ① Initialize  $Q(s, a) \forall s, a$ .  $Q(\text{terminal};) = 0$ .
- ② Repeat for each episode.

Initialize  $s$  using policy derived

Choose  $A \rightarrow$  sampled from  $Q$ .

Repeat (for each step):

Take action  $A$ , get reward  $R$ , landed in  $s'$ .

choose  $A'$  from  $s'$ , using policy derived from  $Q$

$$Q(s, A) \leftarrow Q(s, A) + \alpha (R + \gamma Q(s', A') - Q(s, A))$$

$$s \leftarrow s', A' \leftarrow A$$

Sarsa  $\rightarrow$  on-policy algorithm.

$\downarrow$  converges to the optimal policy, just need a GLIE policy  $\rightarrow$  eg.  $\epsilon_k = 1/k$

- ②  $\alpha$  + step sizes follows.

$$\sum_{t=1}^{\infty} \alpha_t = \infty \quad (\text{sufficiently large to move } Q \text{ values})$$

$$\sum_{t=1}^{\infty} \alpha_t^2 < \infty \quad (\text{step size becomes small})$$

changes to  $\uparrow$  or else noise  
Q values

In practise, sometimes don't depend on both ① & ②. (P)

Again apply similar concept of TD(0)  $\rightarrow$  n-step SARSA.  
to get the best of both worlds.

$n=1 \dots$  (SARSA) —

$n=\infty$  (MC)

n-step  $Q$  return:

$$Q_t^{(n)} \equiv R_{t+1} + \gamma R_{t+2} + \dots + \gamma^n Q(s_{t+n})$$

n-step Q return:

$$Q(s_t, A_t) \leftarrow Q(s_t, A_t) + \alpha (q_t^{(n)} - Q(s_t, A_t))$$

Do this for all ns.  
average over n's

$$\text{Sarsa}(\lambda) \rightarrow Q(s_t, A_t) = Q(s_t, A_t) + \alpha (q_t^\lambda - Q(s_t, A_t))$$

$$q_t^\lambda = (1-\lambda) \sum_{n=1}^{\infty} (\lambda^{n-1}) q_t^{(n)}$$

Forward view.

Backward view  $\rightarrow$  using Eligibility Traces  $\uparrow$

to make algo online

$$E_0(s, a) = 0$$

$$E_t(s, a) = \gamma E_{t-1}(s, a) + I(s_t = s, A_t = a)$$

$Q(s, a) \rightarrow$  updated for every  $(s, a)$  pair

$$\delta_t = R_{t+1} + \gamma Q(s_{t+1}, A_{t+1}) - Q(s_t, A_t)$$

$$Q(s, a) = Q(s, a) + \alpha E_t(s, a) \delta_t$$

• Init  $Q(s, a)$  = arbitrarily

• Repeat - (episode)

$$E(s, a) = 0 \quad \forall s, a.$$

Init  $s, A$ .

Repeat - (step):

Take action  $A$ , observe  $R, s'$

Take  $A' \rightarrow$  from  $Q$ .

$$\delta_t = R_t + \gamma (Q(s', A)) - Q(s, A)$$

$$E(s, A) = E(s, A) + 1$$

$\forall a \neq A$

$$Q(s, a) = Q(s, a) + \alpha \delta_t E(s, a)$$

$$E(s, a) = \gamma E(s, a)$$

$$A \leftarrow A', s \leftarrow s'$$

until termination.

**SARSA( $\lambda$ )**

$\rightarrow$  faster flow of info. back through time

**SARSA(0)**

$\downarrow$   
needs many steps to flow the info. back.



## Off-policy learning

- evaluate target-policy → to compute  $V_\pi / Q_\pi$  following behaviour policy →  $\mu(a/s)$

Why?

- Learn from obsv. of other agents (eg. human) ~~look~~ at traces of behaviour, not just supervised learning

- Re-use experience from old policies  $\pi_1, \pi_2, \dots, \pi_{t-1}$  generated. → using off-policy learning can do that.

- Learn about optimal policy while following exploratory policy

We want policy to be deterministic but also want to explore. optimal off policy learning allows that as we can learn a deterministic optimal policy while following an exploratory policy.

- Learn about multiple policies while following one policy.

Importance sampling → estimate expectation of a different dist

$$E_{X \sim P}[f(X)] = \sum_X P(X) F(X) = \sum_X Q(X) \frac{P(X)}{Q(X)} F(X)$$

eg. reward function

$$= E_{X \sim Q} \left[ \frac{P(X)}{Q(X)} F(X) \right]$$

Use IS for off-policy MC.

- Sample returns from  $\mu$  to evaluate  $\pi$ .
- weigh  $G_T$  as per similarity bet policies
- multiply importance sampling corrections

$$G_T^{\pi_\mu} = \frac{\pi(A_t/s_t)}{\mu(A_t/s_t)} \cdot \frac{\pi(A_{t+1}/s_{t+1})}{\mu(A_{t+1}/s_{t+1})} \dots = \frac{\pi(A_T/s_T)}{\mu(A_T/s_T)} G_T$$



→ Update value towards correct mean.

$$V(s_t) = V(s_t) + \alpha (G_t^{\pi/\mu} - V(s_t))$$

- Very high variance → as multiplying so many ratios diminishes value.
- MC → very bad off-policy

Thus, only TD learning for off-policy

Now, IS only after/up to 1 step.

$$V(s_t) = V(s_t) + \alpha \left( \frac{\pi(A_t/s_t)}{\mu(A_t/s_t)} (R_{t+1} + \gamma V(s_{t+1}) - V(s_t)) \right)$$

- policies need to be similar only over a single step.
- much lower variance than MC (and can still bias off).

Q-learning → Best with off-policy

- make use of Q values.
- no IS reqd.

→ next action from  $\pi$   $A_t \sim \mu(\cdot|s_t)$  ↗ what actually took

→ also an alternative successor action  $A' \sim \bar{\pi}(\cdot|s_t)$   
(what we could have taken following target-policy in future)

- update  $Q(s_t, A_t)$  towards value of successive action
- $$Q(s_t, A_t) \leftarrow Q(s_t, A_t) + \alpha (R_{t+1} + \gamma Q(s_{t+1}, A') - Q(s_t, A_t))$$

main

- allow both behaviour & target policies to improve
- learn greedy policy from exploratory policy.

$\pi \rightarrow$  greedy wrt.  $Q(s, a)$


$$\pi(s_{t+1}) = \underset{a'}{\operatorname{argmax}} Q(s_{t+1}, a')$$

•  $\mu \rightarrow \epsilon$ -greedy w.r.t  $Q(s, a)$ .

Q-learning target:

$$R_{t+1} + \gamma Q(s_{t+1}, A') \\ = R_{t+1} + \gamma \max_{a'} Q(s_{t+1}, a')$$

Q-learning control or Stochastic



~~Q(S, A)~~

$$Q(S, A) \leftarrow Q(S, A) + \alpha (R + \gamma \max_{a'} Q(s', a') - Q(S, A))$$

$\rightarrow$  converges to  $Q^*(S, a)$